

First 90 days:

read deeply, ship narrowly.

A role-appropriate entry plan for building reliable automation capacity inside Vanta Recruiting and People.

Illustrative plan. Sequence and priorities would be adapted to Vanta's existing People Systems architecture, active roadmap, and pod needs.

DAYS 1–30 · READ AND BASELINE

Understand the work before automating it.

Objectives

Build trust with Recruiting, Recruiting Operations, People Operations, People Systems, hiring managers, coordinators, and knowledge owners.

Map the current automation and agent portfolio: what exists, what is adopted, what fails, who owns it, and where knowledge comes from.

Observe real work across sourcing, scheduling, candidate communication, offer-to-hire, reporting, and representative People processes.

Define a lightweight People Automation Warranty and opportunity scorecard.

Core activities

Workflow listening

Architecture reading

Shadow representative pod work; capture manual touches, exceptions, workarounds, candidate/employee friction, decision rights, and existing tool behavior.

Understand Ashby workflows, Dust agents/skills/connectors, Claude use, Zapier, scripts, TeamOhana, Juicebox, LinkedIn, Checkr, Jira, HiBob, Workday, and system boundaries.

Knowledge audit

Portfolio baseline

Map Confluence/Guru sources, owners, freshness, permissions, retrieval gaps, duplication, and missing operating definitions.

Classify workflows by value, manual burden, variability, data readiness, risk, owner strength, adoption burden, reuse, and maintenance load.

Day-30 outputs

Current-state workflow and ownership map.

Automation/agent inventory with keep, repair, observe, or retire hypotheses.

Warranty standard and launch checklist.

Baseline scorecard and opportunity shortlist.

One Recruiting thin slice and one People thin slice selected with accountable owners.

Working agreements to establish

Business ownership

Systems partnership

Every workflow has a Recruiting or People owner who defines the outcome, participates in testing, and owns adoption.

Changes touching architecture, access, core data structure, or platform configuration move with People Systems—not around it.

Human authority

Observable failure

Sensitive candidate and employee decisions remain explicit; automation supports rather than laundering accountability.

Exceptions must be visible, attributable, recoverable, and useful for improving the next release.

What I would not do

Announce a sweeping AI program, replace existing patterns before understanding them, automate an unclear process, or treat tool capability as proof

DAYS 31–60 · BUILD AND VERIFY

Ship thin production slices into real work.

Objectives

Deliver measurable manual-work reduction without hiding exceptions or human authority.

Establish practical Ashby-first patterns and clear boundaries across Dust, Claude, Zapier, scripts, and People Systems.

Instrument workflow behavior, recovery, experience, and adoption from launch.

Representative thin slices

Recruiting slice

People slice

Example: interview scheduling with deterministic availability logic, exception classification, approved communications, coordinator recovery, and time-to-schedule instrumentation.

Example: recurring reporting pipeline with source-linked definitions, scripted validation, AI-assisted summary, visible freshness, and ownerapproved interpretation.

Knowledge slice

Build patterns

Restructure one high-use knowledge domain into owned, retrievable, permission-aware content for Dust agents and People users.

Reusable connectors, skills, logging, error messages, ownership metadata, change tests, and documentation templates.

Verification cadence

Weekly signal

Question

Action

Flow

Did manual touches and cycle time change?

Remove friction or narrow scope.

Reliability

Where did the workflow fail or create rework?

Repair controls and fallback.

Experience

Did recruiters, coordinators, candidates, or employees understand the result?

Fix context, language, or ownership.

Adoption

Are users retaining, bypassing, or modifying the workflow?

Train, redesign, or stop.

Knowledge

Was the agent grounded in current, authorized material?

Repair source or owner.

Adoption plan

Build with users

Teach from the work

Recruiters, coordinators, People partners, and owners shape the thin slice, test exceptions, and approve the operating standard.

Short demos show what changed, why the mechanism fits, what remains human, and how to recover when it fails.

Watch bypass behavior

Close the loop

Edits, workarounds, manual overrides, and abandoned paths become product evidence rather than user resistance to dismiss.

Publish decisions, changes, owners, known limits, and the next review so users know the workflow is maintained.

DAYS 61–90 · STANDARDIZE AND SPREAD

Turn proven workflows into a maintained capability.

Objectives

Scale only patterns that survive real work and earn ownership.

Create an operating cadence across Recruiting, People Operations, and People Systems.

Increase team AI fluency through examples, decision rules, and direct participation.

Make maintenance, knowledge changes, and retirement normal parts of shipping.

Operating system

Portfolio review

Launch standard

Monthly review of value, reliability, adoption, risk, maintenance load, and next decision: scale, repair, human-led, retire.

Warranty, owner, source, action boundary, fallback, instrumentation, documentation, accessibility, security, and change test.

Knowledge protocol

Fluency system

Named owners, freshness expectations, change notifications, permission model, retrieval test, and retirement process.

Build demos, office hours, reusable examples, failure reviews, and champion participation tied to actual People work.

Day-90 readout

Decision

Evidence required

Scale

Reliable performance, clear owner, retained adoption, positive workflow economics, reusable pattern.

Repair

Value signal exists, but source, exception, authority, reliability, or usability remains weak.

Keep human-led

Work is low-volume, high-context, sensitive, or too variable for safe automation.

Retire

Low use, poor value, stale source, high maintenance, duplicated capability, or unacceptable trust burden.

Operating cadence

Weekly build review

Monthly portfolio review

Current slices, exceptions, user feedback, source changes, reliability, and next intervention.

Value, economics, adoption, maintenance burden, reuse, and scale / repair / human-led / retire decisions.

Quarterly architecture review

Failure learning

Connectors, permissions, data structures, knowledge ownership, tool changes, and platform boundaries.

Material failures produce a corrected mechanism, documented reason, and regression check—not a quiet workaround.

90-day success evidence